

Cycle-Consistent Adversarial Networks

Project Report

BY

Name of the Student

Sahdev Nuwal

Varun Banker

Anurag Nagpal

ID Number

2018B3A70900G

2018B2A70295G

2018B1A70939G

Deep Learning CS F425

IC: Prof. Tirtharaj Dash



**BIRLA INSTITUTE OF TECHNOLOGY AND SCIENCE,
PILANI**
Spring semester 2021

Contents

1. Introduction
2. Main Text
 - 2.1. Understanding Cycle-GANS and Architecture
 - 2.2. Python Notebook
 - 2.2.1. Datasets
 - 2.2.2. Training and epochs
 - 2.2.3. Plots
 - 2.3. Results
3. Conclusion and contributions
4. References
5. Thank you!

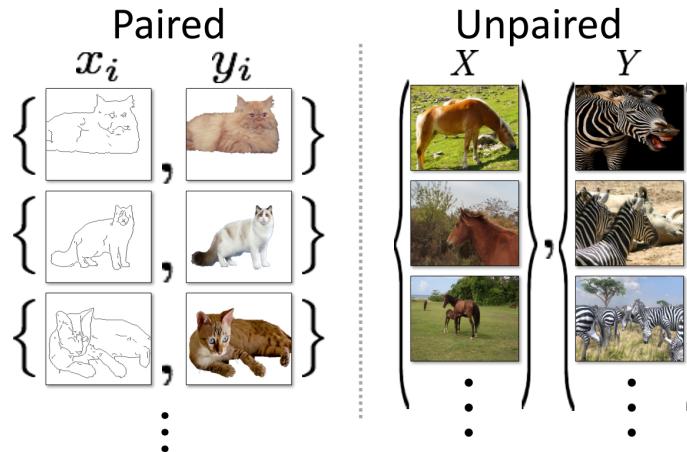
Introduction

Image-to-image translation is a class of vision and graphics problems where the goal is to learn the mapping between an input image and an output image using a training set of aligned image pairs. However, for many tasks, paired training data will not be available. Our project is based on the Paper titled “[Unpaired Image-to-Image Translation using Cycle-Consistent Adversarial Networks](#)” by Jun-Yan Zhu et al 2017. We tried to replicate the architecture of the paper and test the model on three different datasets - [horse2zebra](#), [cityscapes](#), [monet2photo](#).

Main Text

Understanding Cycle GANS and Architecture

CycleGANS results in unpaired image-to-image translation. This means that in order to train it to translate images from domain X to domain Y, we do not have to have exact correspondences between individual images in those domains. The figure below shows Paired training data on the left $\{x_i, y_i\}_{i=1}^N$ where the correspondence between x_i and y_i exists and on the right are unpaired images set, consisting of a source set $\{x_i\}_{i=1}^N$ ($x_i \in X$) and a target set $\{y_j\}_{j=1}^N$ ($y_j \in Y$), with no information provided as to which x_i matches which y_j .



We need to map the input image coordinate X to domain coordinated Y and this problem is an image to image translation. Much literature exists trying to address this problem but every single one of these uses paired image data, their models are trained on both the original image and the corresponding acquired image after translation. But creating such a data paired dataset is difficult and time-consuming. Here lies the motivation behind Cycle GANS, which tries to address this problem using unpaired image to image translation. Unpaired data is much easier to gather.

Suppose for example we have a set of photo style images X and we have another set of Monet's style Y , but we don't have access to Monet's painting for every single input sample image we try to learn a mapping G between images X to its corresponding images Y .

$$G: X \rightarrow Y$$

Our goal is to train a model to learn this mapping G .

Since we don't have paired data, we use cycle consistency loss where we learn two mappings $F: Y \rightarrow X$ and $G: X \rightarrow Y$ and using that we calculate the loss. F and G are inverse mappings of each other. So we not only need a mapping of G that generates similar distribution but we also need one distribution that is cycle consistent with respect to the inverse mapping F .

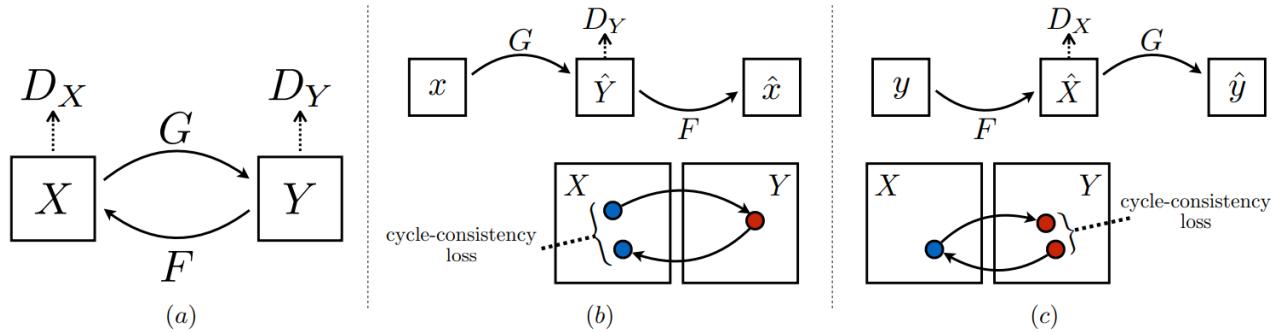
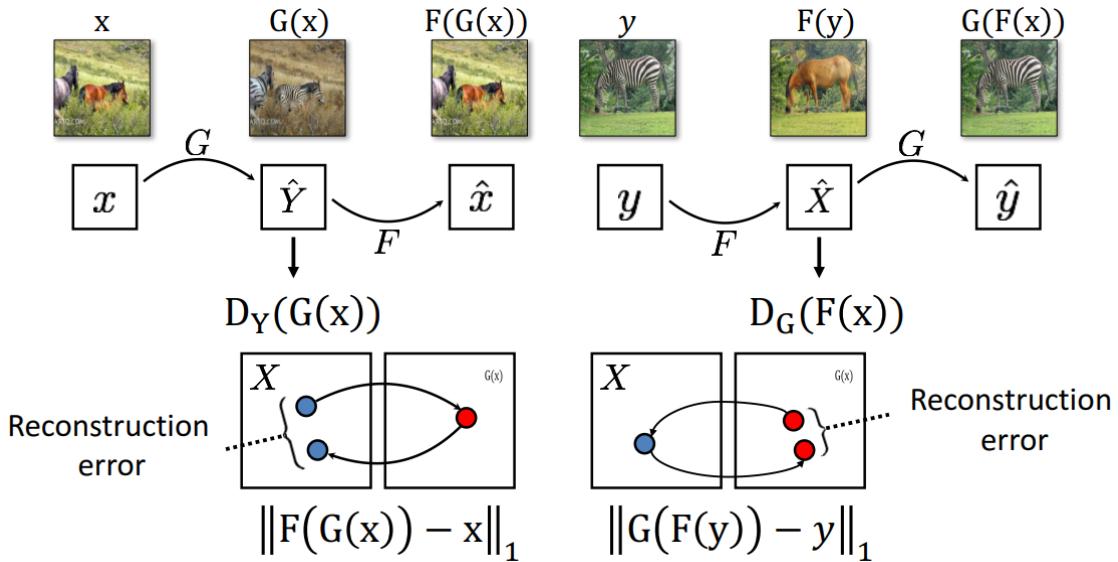


Figure 3: (a) Our model contains two mapping functions $G : X \rightarrow Y$ and $F : Y \rightarrow X$, and associated discriminators D_Y and D_X . D_Y encourages G to translate X into outputs indistinguishable from domain Y , and vice versa for D_X and F . To further regularize the mappings, we introduce two *cycle consistency losses* that capture the intuition that if we translate from one domain to the other and back again we should arrive at where we started: (b) forward cycle-consistency loss: $x \rightarrow G(x) \rightarrow F(G(x)) \approx x$, and (c) backward cycle-consistency loss: $y \rightarrow F(y) \rightarrow G(F(y)) \approx y$

Cycle Consistency Loss



The Reconstruction error is the cycle consistent loss that needs to be minimized. Two kinds of loss functions are used to learn unpaired image to image transformation.

- 1) Adversarial loss: results in style transfer between two domains X and Y

Adversarial Loss: change the style

$$\begin{aligned}\mathcal{L}_{\text{GAN}}(G, D_Y, X, Y) = & \mathbb{E}_{y \sim p_{\text{data}}(y)} [\log D_Y(y)] \\ & + \mathbb{E}_{x \sim p_{\text{data}}(x)} [\log(1 - D_Y(G(x)))]\end{aligned}$$

Similarly, for F , we also have $\mathcal{L}_{\text{GAN}}(F, D_X, X, Y)$.

- 2) Cycle Consistency Loss (Reconstruction Error) Which preserves the content

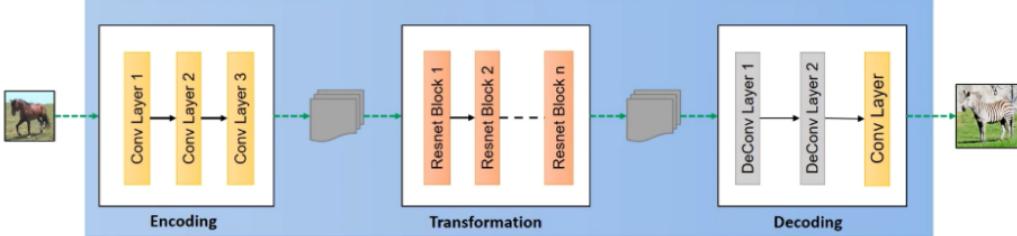
Cycle Consistency Loss: preserve the content

$$\begin{aligned}\mathcal{L}_{\text{cyc}}(G, F) = & \mathbb{E}_{x \sim p_{\text{data}}(x)} [\|F(G(x)) - x\|_1] \\ & + \mathbb{E}_{y \sim p_{\text{data}}(y)} [\|G(F(y)) - y\|_1].\end{aligned}$$

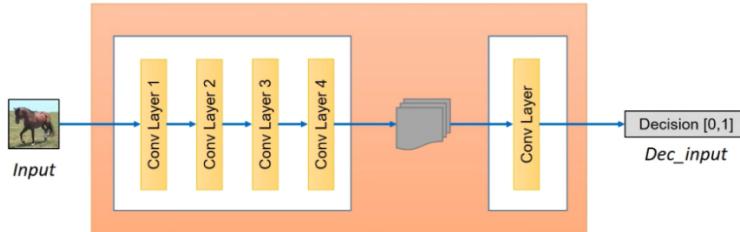
Two empirical assumptions are that style is easy to change and content is easy to keep.

Therefore, the full objective is:

$$\begin{aligned}\mathcal{L}(G, F, D_X, D_Y) = & \mathcal{L}_{\text{GAN}}(G, D_Y, X, Y) \\ & + \mathcal{L}_{\text{GAN}}(F, D_X, Y, X) \\ & + \lambda \mathcal{L}_{\text{cyc}}(G, F)\end{aligned}$$



Generator Architecture



Discriminator Architecture

Algorithm 2 CycleGAN Training Loop Pseudocode

- 1: **procedure** TRAINCYCLEGAN
- 2: Draw a minibatch of samples $\{x^{(1)}, \dots, x^{(m)}\}$ from domain X
- 3: Draw a minibatch of samples $\{y^{(1)}, \dots, y^{(n)}\}$ from domain Y
- 4: Compute the discriminator loss on real images:

$$\mathcal{J}_{real}^{(D)} = \frac{1}{m} \sum_{i=1}^m (D_X(x^{(i)}) - 1)^2 + \frac{1}{n} \sum_{j=1}^n (D_Y(y^{(j)}) - 1)^2$$

- 5: Compute the discriminator loss on fake images:

$$\mathcal{J}_{fake}^{(D)} = \frac{1}{m} \sum_{i=1}^m (D_Y(G_{X \rightarrow Y}(x^{(i)})))^2 + \frac{1}{n} \sum_{j=1}^n (D_X(G_{Y \rightarrow X}(y^{(j)})))^2$$

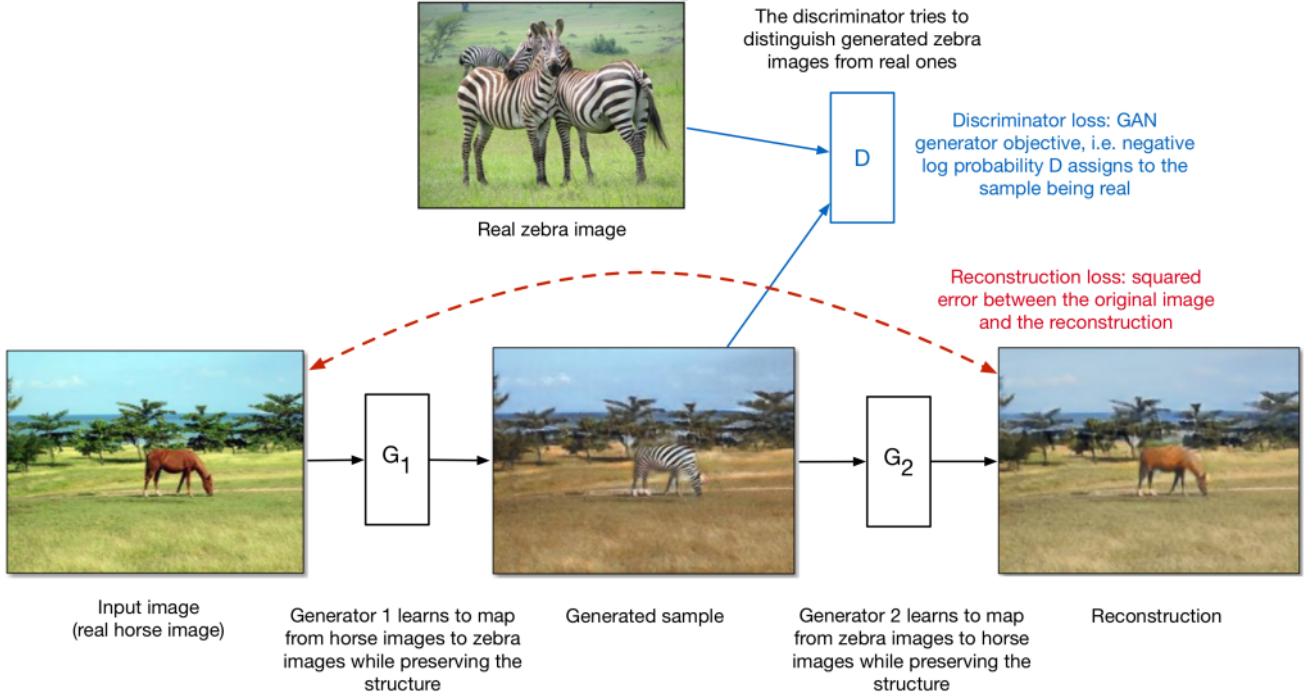
- 6: Update the discriminators
- 7: Compute the $Y \rightarrow X$ generator loss:

$$\mathcal{J}^{(G_{Y \rightarrow X})} = \frac{1}{n} \sum_{j=1}^n (D_X(G_{Y \rightarrow X}(y^{(j)})) - 1)^2 + \mathcal{J}_{cycle}^{(Y \rightarrow X \rightarrow Y)}$$

- 8: Compute the $X \rightarrow Y$ generator loss:

$$\mathcal{J}^{(G_{X \rightarrow Y})} = \frac{1}{m} \sum_{i=1}^m (D_Y(G_{X \rightarrow Y}(x^{(i)})) - 1)^2 + \mathcal{J}_{cycle}^{(X \rightarrow Y \rightarrow X)}$$

- 9: Update the generators
-



Datasets

We used three different datasets from those used in the paper and one new but similar dataset. The three data sets are [horse2zebra](#), [cityscapes](#), and [monet2photos](#).

Training and Epochs

Total number of epochs = 29

Number of Iterations per epoch = 214

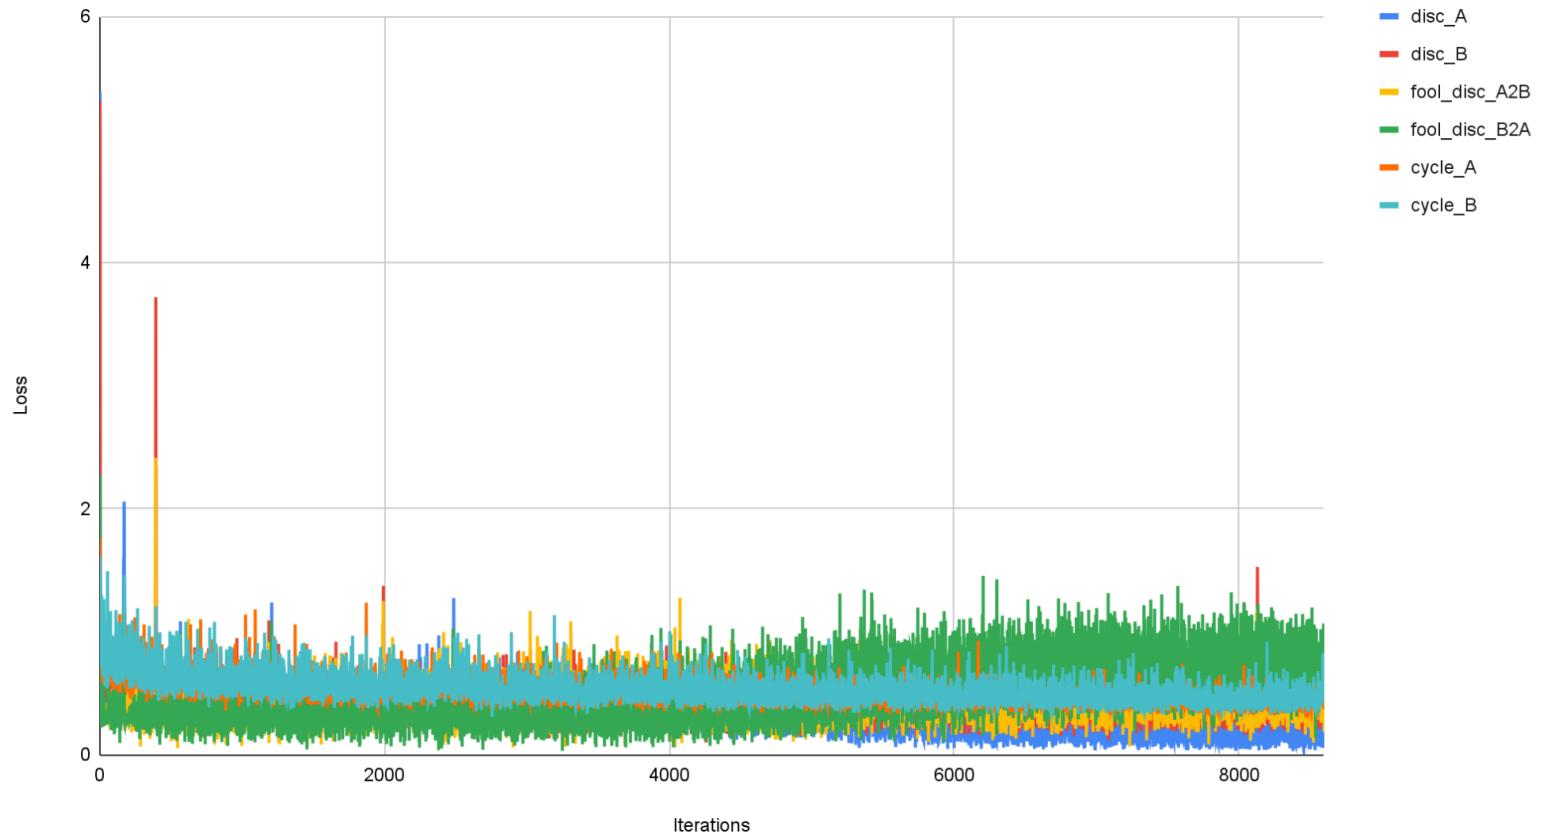
Adam Optimizer was used

Add screenshots of last training epochs

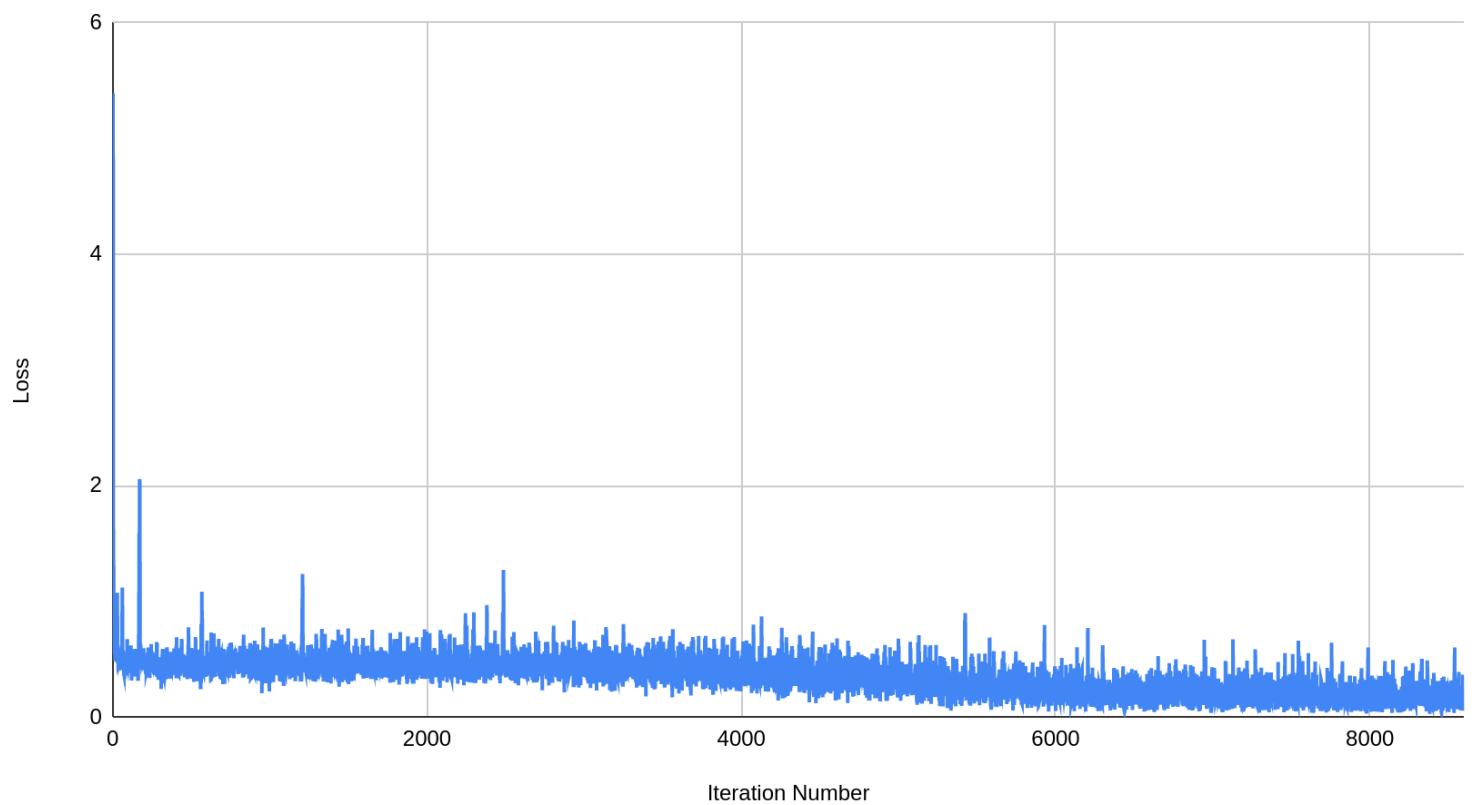
```
Debug Output:  
Iteration: 720  
Disc A Loss: 0.06755250692367554  
Disc B Loss: 0.3581489622592926  
Fool Disc Loss A2B: 0.28648442029953003  
Fool Disc Loss B2A: 0.8311122059822083  
Cycle Loss A: 0.4563145041465759  
Cycle Loss B: 0.43075913190841675  
  
Debug Output:  
Iteration: 721  
Disc A Loss: 0.08424031734466553  
Disc B Loss: 0.3619263768196106  
Fool Disc Loss A2B: 0.5194675326347351  
Fool Disc Loss B2A: 1.0486863851547241  
Cycle Loss A: 0.4273029565811157  
Cycle Loss B: 0.36301401257514954  
  
Debug Output:  
Iteration: 722  
Disc A Loss: 0.16360032558441162  
Disc B Loss: 0.3867865800857544  
Fool Disc Loss A2B: 0.5702834725379944  
Fool Disc Loss B2A: 0.7987803816795349  
Cycle Loss A: 0.37391197681427  
Cycle Loss B: 0.3457791209220886  
  
Debug Output:  
Iteration: 723  
Disc A Loss: 0.12710599601268768  
Disc B Loss: 0.2555253505706787  
Fool Disc Loss A2B: 0.49111172556877136  
Fool Disc Loss B2A: 0.4880658686161041  
Cycle Loss A: 0.3765408396720886  
Cycle Loss B: 0.3894083499908447  
  
Debug Output:  
Iteration: 724  
Disc A Loss: 0.21715083718299866  
Disc B Loss: 0.27839767932891846  
Fool Disc Loss A2B: 0.495213121175766  
Fool Disc Loss B2A: 0.8631212711334229  
Cycle Loss A: 0.3724614083766937  
Cycle Loss B: 0.4944200813770294
```

Plots

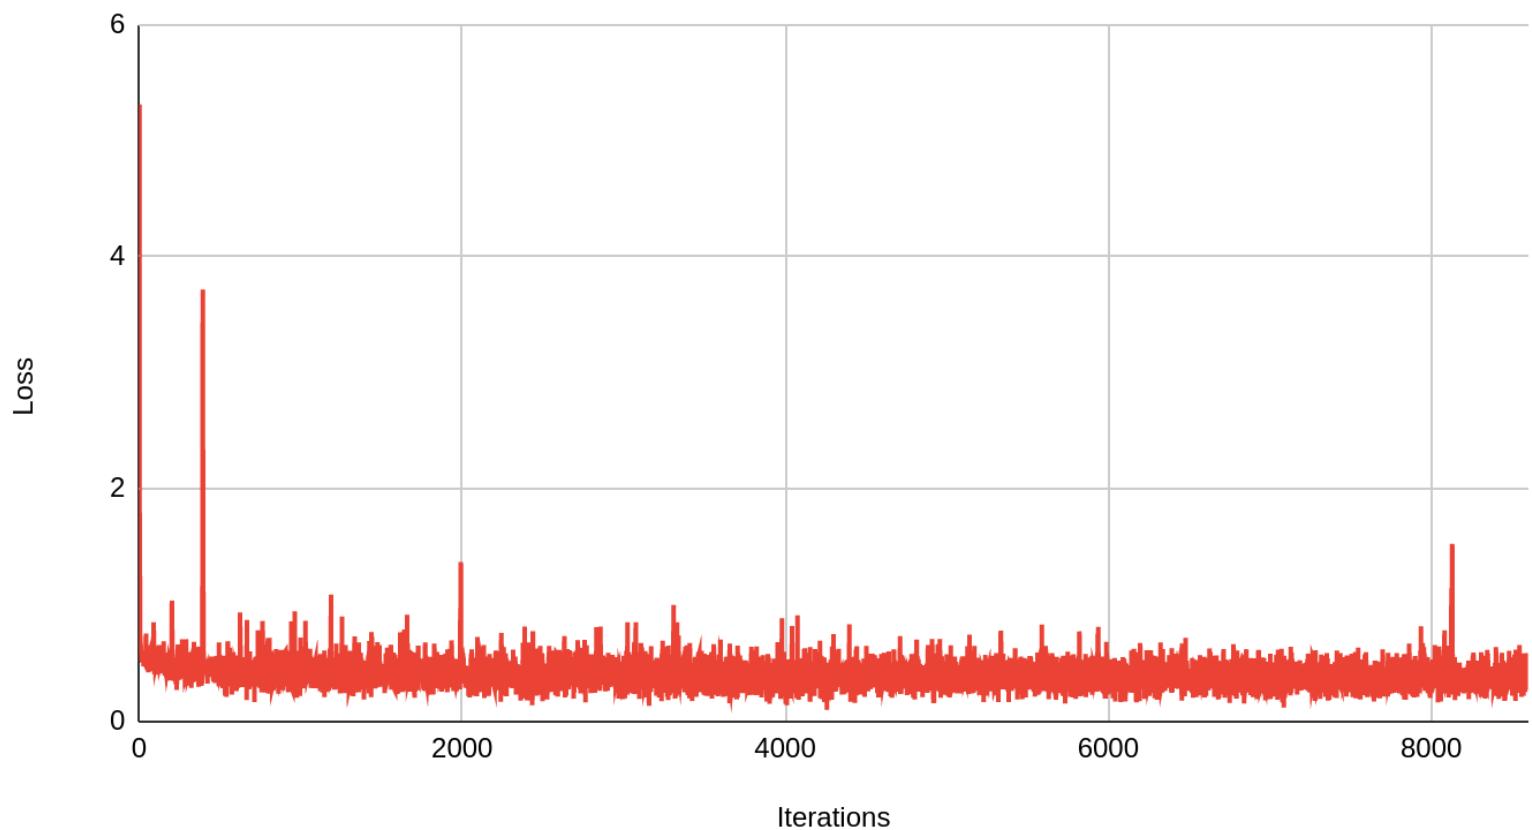
All losses combined



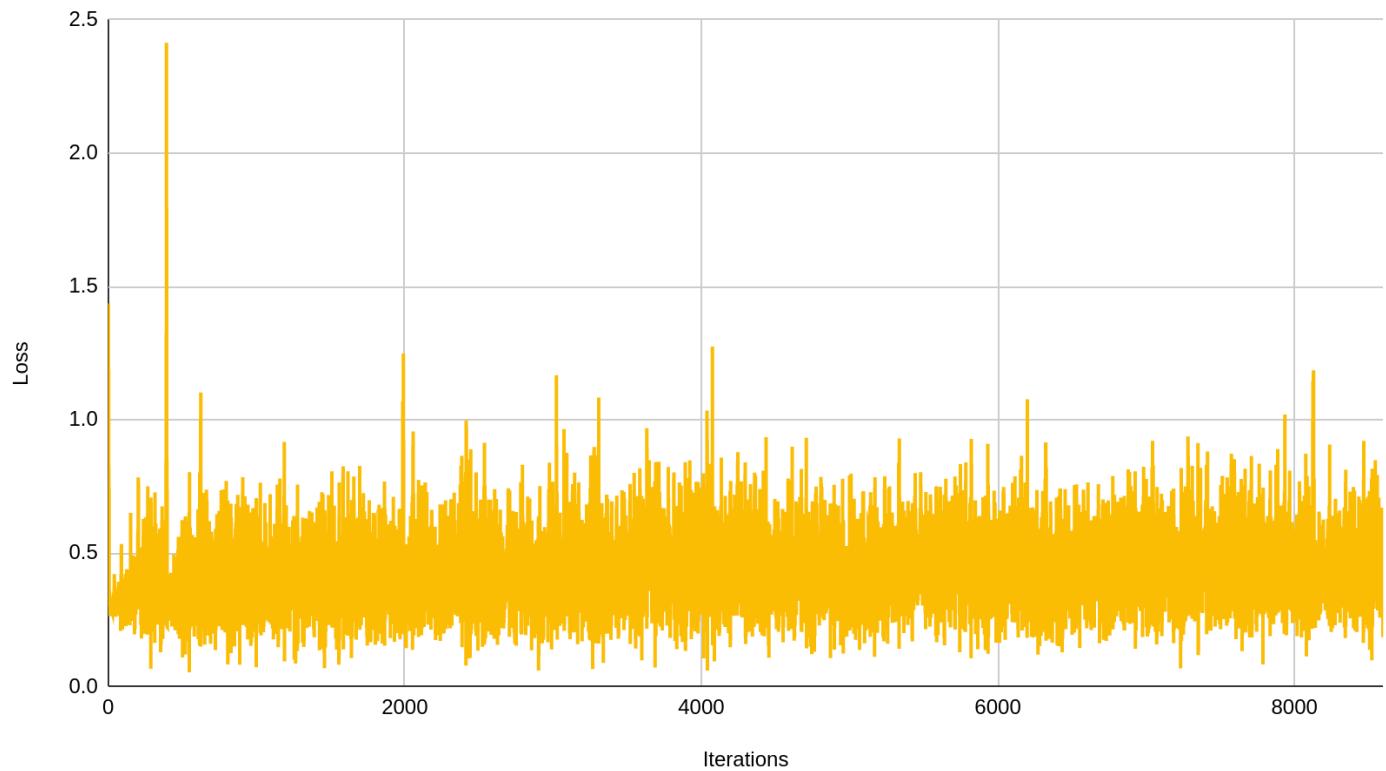
Disc_A



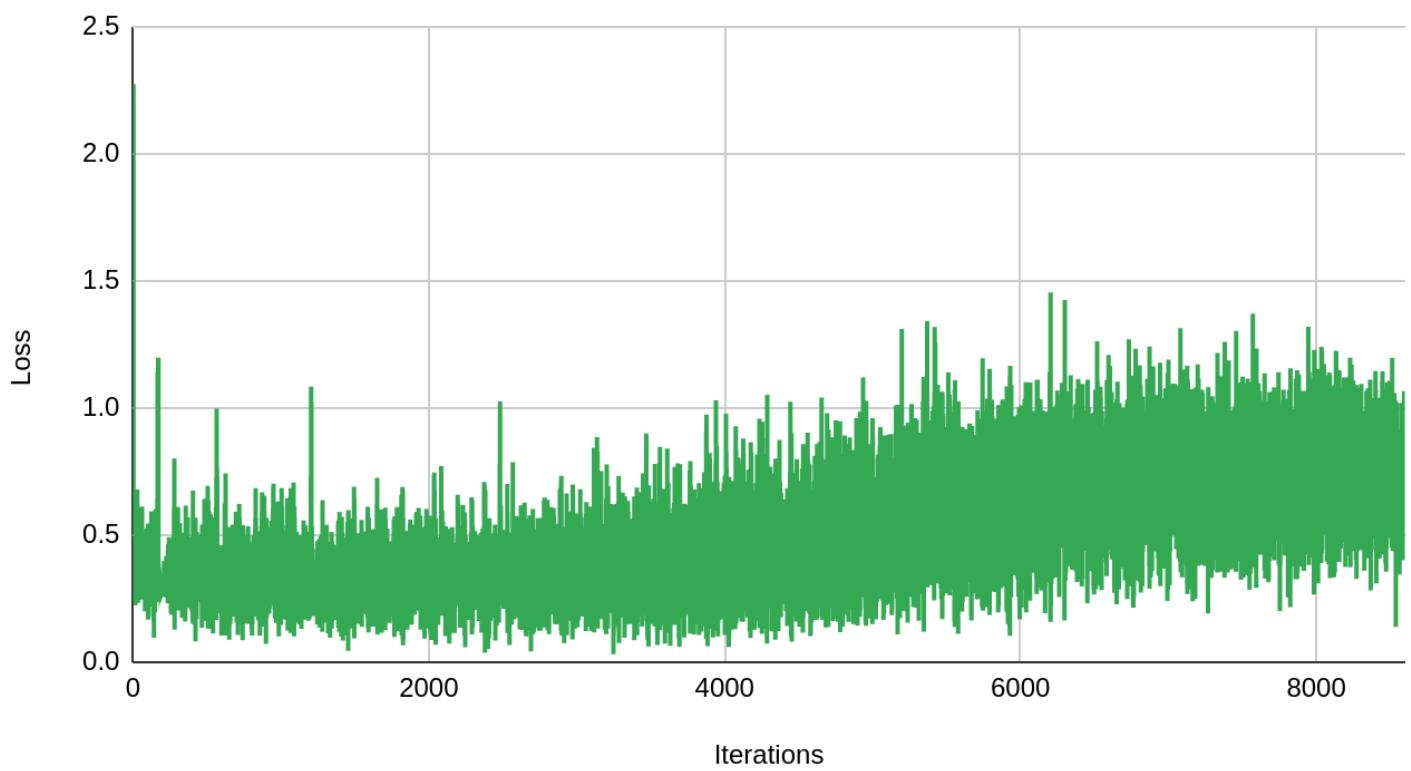
Disc_B



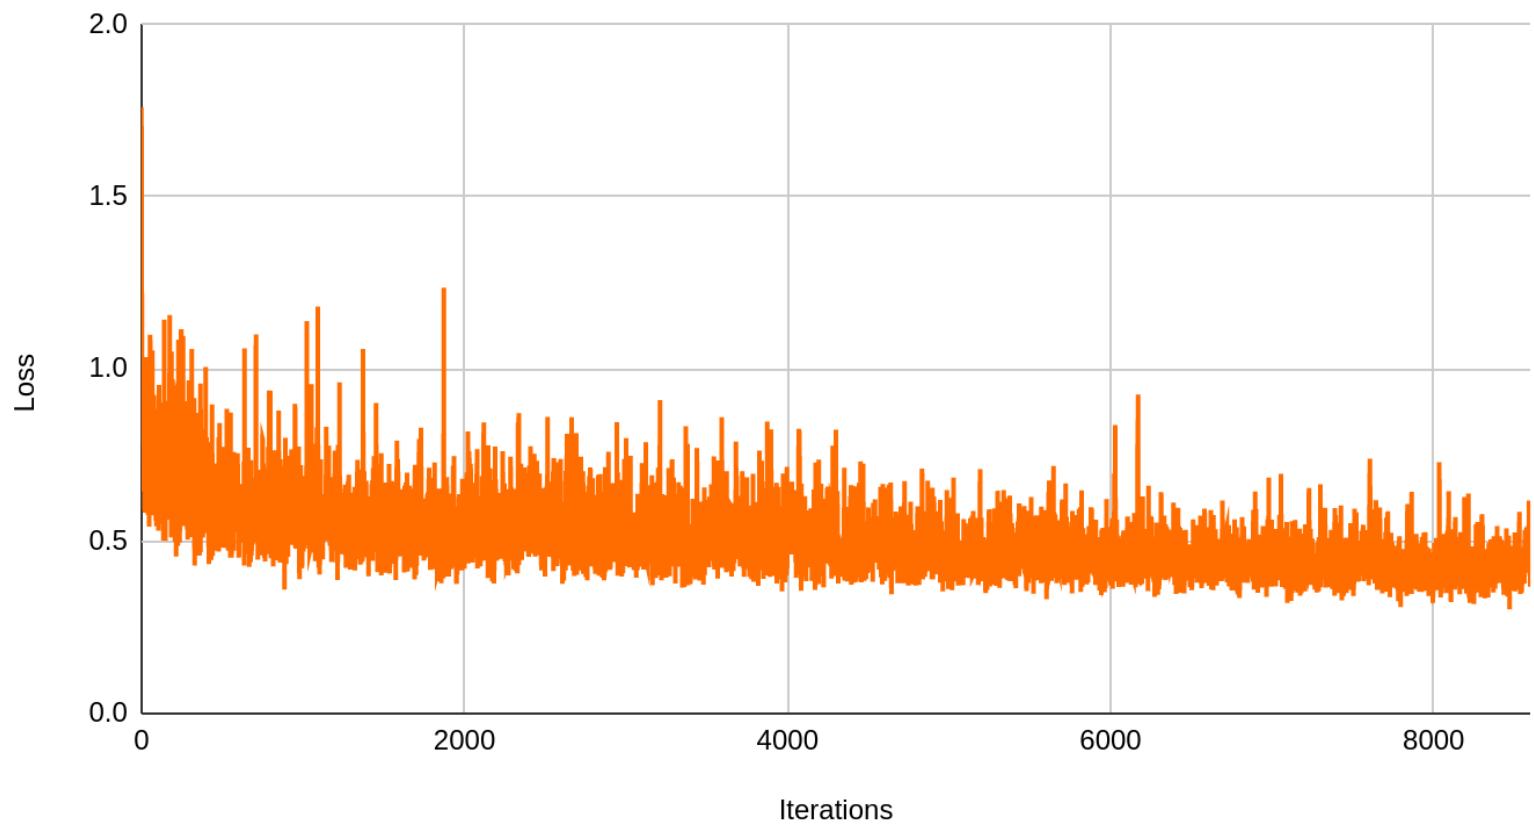
Fool_disc_A2B



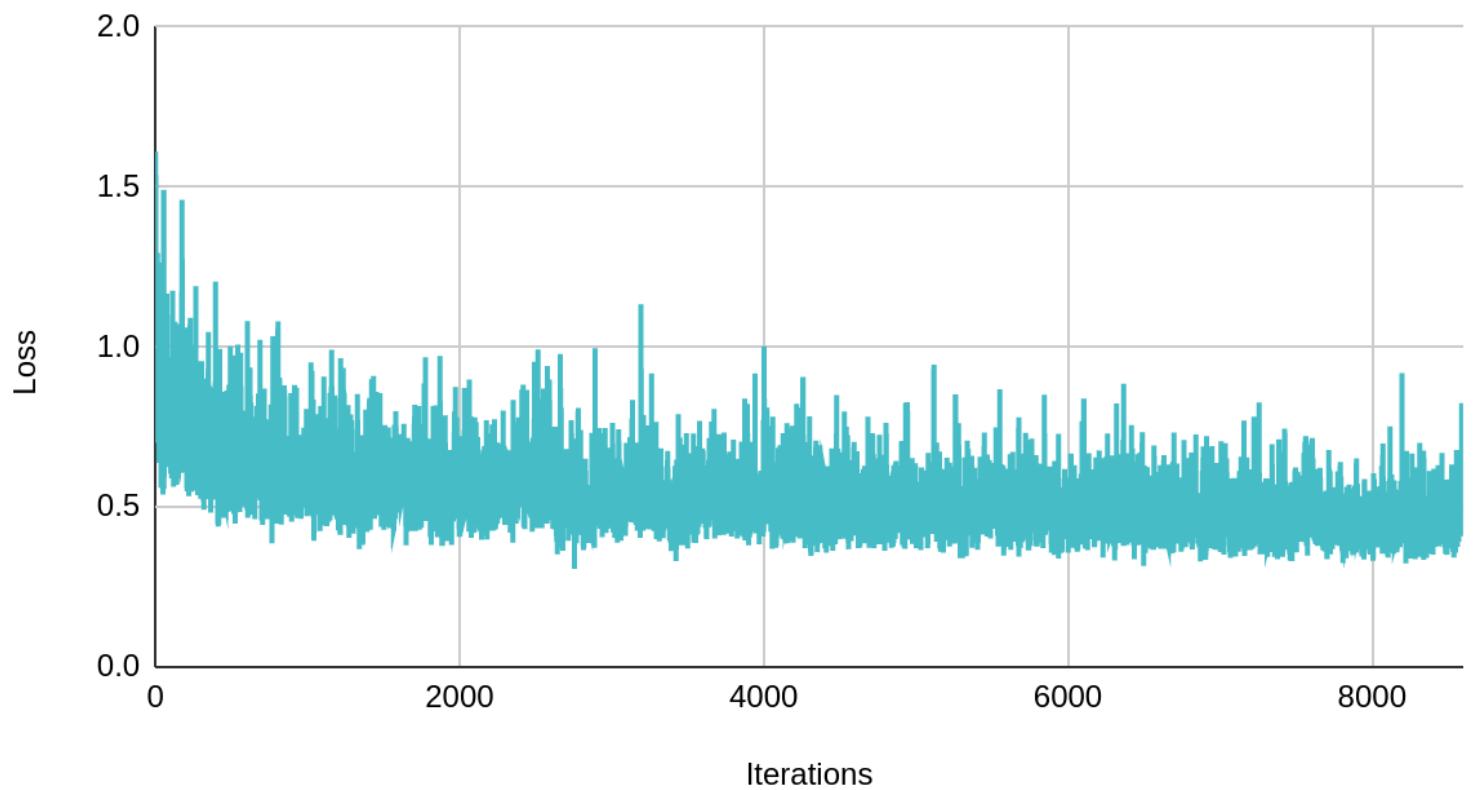
Fool_disc_B2A



Cycle_A



Cycle_B



Results

Real monet



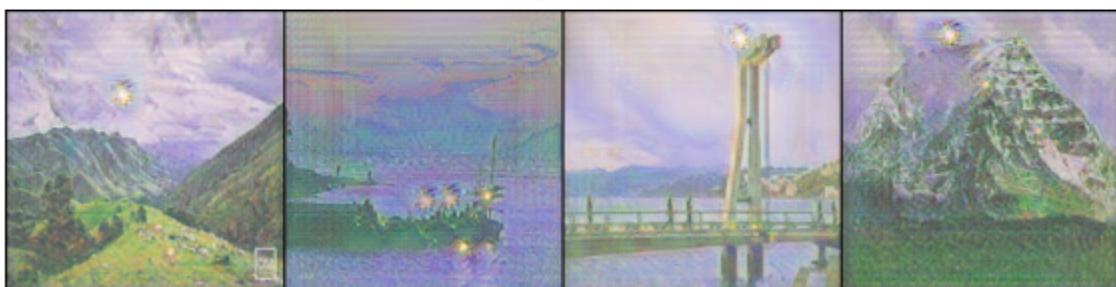
Fake photo



Real photo



Fake monet



Identity monet



Identity photo



Recovered monet



Recovered photo



Conclusion and contributions

In conclusion we learned that Cycle Consistency Adversarial Networks tries to solve image-image translation problems without paired data. The objective function it uses is the Adversarial losses and Cycle consistency losses. Image-Image translation is used in Object Transfiguration, Photo Enhancement, Style Transformation, Season Transformation and Image Segmentation. We made the model and tested it on a new city dataset and performed image segmentation.

References

- 1) Zhu, Jun-Yan, et al. "Unpaired image-to-image translation using cycle-consistent adversarial networks." *Proceedings of the IEEE international conference on computer vision*. 2017.
- 2) [Jun-Yan Zhu and Taesung Park Joint work with Phillip Isola and Alexei A. Efros](#)
- 3) [CSC321 Lecture 19: Generative Adversarial Networks](#)
- 4) [CycleGAN Introduction Important Note](#)